

# Computer Science Fundamentals

Number Systems – Binary Multiplication & Division, Floating-Point Numbers

Technische Hochschule Rosenheim  
Winter 2021/22  
Prof. Dr. Jochen Schmidt

- Binary multiplication, division
- Floating-point numbers

- Rules for the multiplication of two binary digits
  - $0 \cdot 0 = 0$
  - $0 \cdot 1 = 0$
  - $1 \cdot 0 = 0$
  - $1 \cdot 1 = 1$
- Identical to the rules of logical AND!
- Multiplication of multi-digit numbers
  - Multiplication of the multiplicand by the individual digits of the multiplier
  - Proper addition (at the correct position) of the interim results

# Binary Multiplication – Example

Example:  $10 \cdot 13$

$$\begin{array}{r} 1010 \cdot 1101 \\ \hline 1010 \\ 10100 \\ 00000 \\ 101000 \\ \hline 10000010 \end{array}$$

**Result:**  $130_{10}$

## Typically

- with fixed number of bits
- multiplier digits from right to left (instead of left to right as in the previous slide)
- shift operations

Example: 8 Bits,  $10 \cdot 13 = 0000\ 1010 \cdot 0000\ 1101$

0000 1010	1
0000 0000	0, shift multiplicand left by 1 bit
0010 1000	1, shift left by 2 bits
0101 0000	1, shift left by 3 bits
0000 0000	0, shift left by 4 bits
...	
<u>0000 0000</u>	0, shift left by 7 bits
1000 0010	= 130

Note:

We do not have to store store all these and add at the end.

We can directly add each shifted multiplicand and get an intermediate result

# Binary Multiplication – Example

This also works in **complement** representation with negative numbers (example: 8 Bits, two's complement)

$$(-5) \cdot 3 = 1111\ 1011 \cdot 0000\ 0011$$

1111 1011	1
1111 0110	1, shift left by 1 bit
0000 0000	0, shift left by 2 bits
...	
0000 0000	0, shift left by 7 bits
<u>101111 0001</u>	= -15
overflow discarded	

$$3 \cdot (-5) = 0000\ 0011 \cdot 1111\ 1011$$

0000 0011	1
0000 0110	1, shift left by 1 bit
0000 0000	0, shift left by 2 bits
0001 1000	1, shift left by 3 bits
0011 0000	1, shift left by 4 bits
0110 0000	1, shift left by 5 bits
1100 0000	1, shift left by 6 bits
<u>11000 0000</u>	1, shift left by 7 bits
101111 0001	= -15
overflows discarded	

$$(-3) \cdot (-5) = 1111\ 1101 \cdot 1111\ 1011$$

1111 1101	1
1111 1010	1, shift left by 1 bit
0000 0000	0, shift left by 2 bits
1110 1000	1, shift left by 3 bits
1101 0000	1, shift left by 4 bits
1010 0000	1, shift left by 5 bits
0100 0000	1, shift left by 6 bits
<u>1000 0000</u>	1, shift left by 7 bits
0000 1111	= +15
overflows discarded	

note: there are also other ways to do this multiplication

# Binary Multiplication – Fixed-Point Example

Example:  $17.375 \cdot 9.75$

$$\begin{array}{r} 1\ 0\ 0\ 0\ 1.0\ 1\ 1\ .\ 1\ 0\ 0\ 1.1\ 1 \\ \hline \phantom{1\ 0\ 0\ 0\ 1.0\ 1\ 1\ .\ 1\ 0\ 0\ 1.1\ 1} 1\ 0\ 0\ 0\ 1\ 0\ 1\ 1 \\ \phantom{1\ 0\ 0\ 0\ 1.0\ 1\ 1\ .\ 1\ 0\ 0\ 1.1\ 1} 1\ 0\ 0\ 0\ 1\ 0\ 1\ 1 \\ \phantom{1\ 0\ 0\ 0\ 1.0\ 1\ 1\ .\ 1\ 0\ 0\ 1.1\ 1} 1\ 0\ 0\ 0\ 1\ 0\ 1\ 1 \\ \phantom{1\ 0\ 0\ 0\ 1.0\ 1\ 1\ .\ 1\ 0\ 0\ 1.1\ 1} 1\ 0\ 0\ 0\ 1\ 0\ 1\ 1 \\ \phantom{1\ 0\ 0\ 0\ 1.0\ 1\ 1\ .\ 1\ 0\ 0\ 1.1\ 1} 1\ 0\ 0\ 0\ 1\ 0\ 1\ 1 \\ \hline \text{Intermediate result} \phantom{1\ 0\ 0\ 0\ 1.0\ 1\ 1\ .\ 1\ 0\ 0\ 1.1\ 1} 1\ 0\ 1\ 0\ 1\ 0\ 0\ 1.0\ 1\ 1\ 0\ 1 \end{array}$$

Insert point at correct position

**Result:**  $169.40625_{10}$

Basically the same as decimal

Example 20 : 6

$$\begin{array}{r} 10100 : 110 = 11.0101 \dots \\ \underline{-110} \phantom{000} \\ 1000 \phantom{00} \\ \underline{-110} \phantom{00} \\ 1000 \phantom{00} \\ \underline{-110} \phantom{00} \\ 1000 \phantom{00} \\ \underline{-110} \phantom{00} \\ \dots \end{array}$$


**Result: 3.333...<sub>10</sub>**



If multiplier or divisor is a **power of two**  $2^k$

- Multiplication or division can be done easier and faster
- By shifting a corresponding number of bits ( $k$ ) to the left or right
- For  $2^1$  by 1 Bit, for  $2^2$  by 2 bits, for  $2^3$  by 3 bits etc.

# Special Case – Shift – Examples



- $13 \cdot 4$

$$1101 \cdot 100 = 110100$$

- $20 \cdot 8$

$$10100 \cdot 1000 = 10100\textcolor{brown}{000}$$

← 3 bits to the left

- $20 : 4$

$$101\textcolor{brown}{00} : 100 = 101$$

→ 2 bits to the right

- $26 : 4$

$$110\textcolor{brown}{10} : 100 = 110 \text{ (Remainder } \textcolor{brown}{2} \text{)} \quad (\rightarrow \text{possible loss of information!})$$

→ 2 bits to the right

				1	1	0	1
			1	1	0	1	0
		1	1	0	1	0	0

2 bits to the left

= 52

- we cannot represent real numbers in a computer
  - we'll always have to use a finite number of bits
  - so, we actually have rational numbers only, represented as (binary) fractions
- Two main types
  - fixed-point numbers (*Festkommazahlen*)
  - floating-point numbers (*Gleitkommazahlen*)

- Point separating integer from fractional part always at the same position

- therefore, the point itself does not have to be stored

- $Z_2$  has length  $n + m$  bit

**n** digits left and **m** digits right of the point

$$Z_2 = z_{n-1}z_{n-2} \cdots z_1z_0 \cdot z_{-1}z_{-2} \cdots z_{-m}$$

$$Z_{10} = \sum_{i=-m}^{n-1} z_i \cdot 2^i$$

- Disadvantages of fixed-point arithmetic

- Using a fixed number of bits, only a limited range of values can be covered
  - The location of the point is always the same

(Where to put it, if sometimes you have to operate with very small, highly accurate values, and at other times with very large values?)

- Fixed-point arithmetic is only used in special purpose computers

- otherwise: floating-point arithmetic

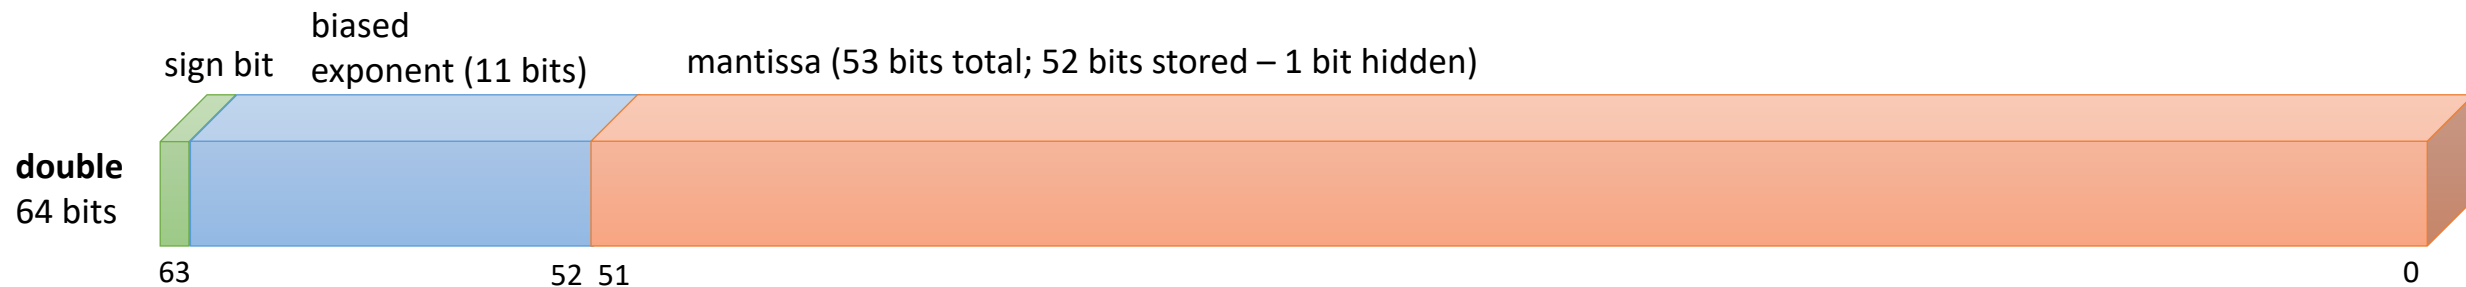

- Example: Any decimal fraction can be written in the following form

$$2.3756 \cdot 10^3$$

- Two components
  - **Mantissa (or Signifand)** (2.3756) and
  - **Exponent** (3), which is integer
- Used in most computers
  - with base 2 instead of base 10
- We have to specify, how many bits to use for the representation; most common:
  - **single precision** (32 bits, data type **float**) or
  - **double precision** (64 bits, data type **double**)



- standard IEEE 754-2019
- C/C++ and Java data types use 4 bytes for **float** and 8 bytes for **double**
- the standard also defined types for **half** (16 Bit) and **quadruple** (128 Bit) precision



- we use **normalized** floating-point numbers
  - the exponent is changed
  - such that the (not stored) point is always directly to the right of the first non-zero digit
  - which, in binary is always 1  $\rightarrow$  no need to store it (the hidden bit of the mantissa)

- Example:  $17.625_{10}$

$$= 16 + 1 + \frac{1}{2} + \frac{1}{8}$$

$$= 10001.101_2$$

$$= 10001.101 \cdot 2^0 \quad \text{this notation is a bit of a mixture binary – decimal}$$

- Normalized form
  - move the point directly to the right of the first significant digit
  - change the exponent accordingly

$$= 1.0001101 \cdot 2^4$$

- In the **mantissa** the most significant bit to the left of the (not stored) point is always 1
  - → no need to store it (the hidden bit of the mantissa)
  - except for 0.0 and some other special cases
- **Exponent** is an integer, which (after adding a **bias**) can be represented without a sign
  - the value of the **bias** depends on precision
    - float (4 bytes, 8 bits for exponent): Bias = 127
    - double (8 bytes, 11 bits for exponent): Bias = 1023
  - using **bias** addition, no special sign treatment is required for exponent arithmetic (always positive)
  - it would have been possible to use complement representation instead
    - but bias representation makes comparison of floating-point numbers easier
- **Sign bit** indicates the sign of the mantissa
  - mantissa stored as absolute value
  - sign bit = 0 → positive
  - sign bit = 1 → negative



# IEEE Floating-point Format – Example

- 17.625 ( $1.0001101 \cdot 2^4$ )

3 1	3 0	2 9	2 8	2 7	2 6	2 5	2 4	2 3	2 2	2 1	2 0	1 9	1 8	1 7	1 6	1 5	1 4	1 3	1 2	1 1	1 0	9	8	7	6	5	4	3	2	1	0
0	1	0	0	0	0	0	1	1	0	0	0	1	1	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

- Biased exponent: 10000011 = 131
- Bias: 01111111 = 127
- Real exponent: 00000100 = 4

Precision	Single	Double
Sign bits	1	1
exponent bits	8	11
mantissa bits	23	52
bits total	32	64
Bias	127	1023
exponenten range	[-126, 127]	[-1022, 1023]

- positive (+0.0) and negative (−0.0) zero
  - these compare to „equal“!
  - (biased) exponent = 0
  - mantissa = 0
  - sign = 0/1
- Usage
  - Representation of zero
  - Rounding to  $\pm 0.0$  for underflows (there's a gap around zero – *underflow gap*)

- plus ( $+\infty$ ) and minus ( $-\infty$ ) infinity
  - (biased) exponent = 111....1
  - mantissa = 0
  - sign = 0/1
- Usage
  - Numbers with too large absolute values to be represented (overflow)
  - Computations that by definition result in infinity  
(e.g., division of a number  $z \neq 0$  by zero:  $z / 0.0 = \infty$ )

- NaN: Not a Number
  - (biased) exponent = 111....1
  - mantissa > 0
  - sign = 0/1
- Usage
  - Representation of invalid values
  - Computations that provide undefined results, e.g.,
    - $0.0 / 0.0$  = NaN
    - $\infty / \infty$  = NaN
    - $\sqrt{-3}$  = NaN
- comparisons with NaNs always result in *false*
  - even when testing  $\text{NaN} == \text{NaN}$  ( $\rightarrow$  *false*)

Convert the decimal number 125.875 to `float` ( IEEE format)

1. Convert to binary as usual, ignore sign

Integer part:  $125_{10} = 1111101_2$

Fractional part:  $0.875_{10} = 0.111_2$

$$0.875 \cdot 2 = 1.75 \rightarrow 1$$

$$0.75 \cdot 2 = 1.5 \rightarrow 1$$

$$0.5 \cdot 2 = 1.0 \rightarrow 1$$

2. Normalize

$$1111101.111 \cdot 2^0 = 1.111101111 \cdot 2^6$$

3. Determine exponent in binary

bias for float:  $127_{10}$

$$2^6 \rightarrow 6_{10} + \text{bias} = 133_{10} = 10000101_2$$

4. Determine sign bit: positive  $\rightarrow 0$

5. Combine results

s	exponent	mantissa
0	10000101	1111011110000000000000

- Floating-point numbers that can be accurately represented in the decimal system cannot always be accurately represented in the dual system
  - Note: Never compare `float` or `double` values for equality!
- Examples on the following slides: Print numbers from 0.1 to 1.0 using step 0.1

Loop using floats as counters: infinite loop, termination condition is never reached

```
#include <stdio.h>

int  main(void)
{
    float  i = 0.1;

    while (i != 1.0)
    {
        printf("%.10f\n", i);
        i = i + 0.1;
    }
    return 0;
}
```



This gives the desired result:

```
#include <stdio.h>

const float EPSILON = 1e-6;

int main(void)
{
    float i = 0.1;

    while (i <= 1.0+EPSILON)
    {
        printf("%.10f\n", i);
        i = i + 0.1;
    }
    return 0;
}
```

Better: Avoid inaccuracies by using integers as loop counters

```
#include <stdio.h>

int main(void) {
    int i = 1;

    while (i <= 10)
    {
        printf("%.10f\n", (float)i/10);
        i = i + 1;
    }
    return 0;
}
```

- Floating-point arithmetic is not associative!

- $(u + v) + w \neq u + (v + w)$   
 $(u \cdot v) \cdot w \neq u \cdot (v \cdot w)$

- ...and not distributive

- $u \cdot (v + w) \neq (u \cdot v) + (u \cdot w)$

- Example (decimal, accuracy of 8 digits)

- $(11111113. + (-11111111.)) + 7.5111111 =$   
 $2.0000000 + 7.5111111 = 9.5111111$

- $11111113. + (-11111111. + 7.5111111) =$   
 $11111113. + (-11111103.) = 10.0000000$